

dav

October 17, 2024

```
[ ]: # TODO:
# 1.      Import
# 2.      Read csv
# 3.      Head
# 4.      Shape
# 5.      Describe
# 6.      Isnull().sum()
# 7.      Dropna(inplace=True)
# 8.      Isnull().sum()
# 9.      Separate numerical and categorical cols using select_dtypes
# 10.     Calculate mean, median, std dev, quartiles
# 11.     Correlation matrix (heatmap)
# 12.     Find top 5 features from heatmap
# 13.     Plot Histogram (all numerical cols)
# 14.     Box Plot (all numerical cols)
# 15.     Scatter Plot (all numerical cols, upon a target)
# 16.     Feature Engineering
# 17.     Define function to find outliers (IQR) Note: if df.quantile(float_
    ↪ val); np.percentile(whole val)
# 18.     Find number of outliers
# 19.     Remove outliers or fill with median
# 20.     Display bar chart for categorical variable
# 21.     Display mean, median, mode, std dev, IQR
# 22.     Scale the numerical data - MinMax or StandardScaler
# 23.     Select one column for KDE histogram, making it normal distribution
# 24.     Plot regular histogram and KDE histogram
# 25.     Plot QQ plot (stats.probplot)
# 26.     Do Shapiro and KSTest (stats.shapiro and stats.kstest)
# 27.     If p values are > 0.05, normal or else not normal dist
# 28.     Transformation using np.log1p or boxcox (stats.boxcox)
# 29.     Histogram and QQ Plot for new transformed value
# 30.     Do those tests again and check if p > 0.05
# 31.     Select independent and dependent variable (for linear regression)
    ↪ (X - [])
# 32.     Split the data for train and test (0.2 for test size)
# 33.     Create model and fit
```

```
# 34. Find intercept and slope (intercept -> model.intercept_, slope ->
    ↪model.coef_[0]) and print regression eqn
# 35. Predict using model (X_test)
# 36. Plot regression line
# 37. Display metrics
# 38. Display residual graph
```

```
[47]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as stats
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
[48]: url = "https://raw.githubusercontent.com/ageron/handson-ml/refs/heads/master/
    ↪datasets/housing/housing.csv"
df = pd.read_csv(url)
```

```
[49]: df.head()
```

```
[49]: longitude latitude housing_median_age total_rooms total_bedrooms \
0 -122.23 37.88 41.0 880.0 129.0
1 -122.22 37.86 21.0 7099.0 1106.0
2 -122.24 37.85 52.0 1467.0 190.0
3 -122.25 37.85 52.0 1274.0 235.0
4 -122.25 37.85 52.0 1627.0 280.0

population households median_income median_house_value ocean_proximity
0 322.0 126.0 8.3252 452600.0 NEAR BAY
1 2401.0 1138.0 8.3014 358500.0 NEAR BAY
2 496.0 177.0 7.2574 352100.0 NEAR BAY
3 558.0 219.0 5.6431 341300.0 NEAR BAY
4 565.0 259.0 3.8462 342200.0 NEAR BAY
```

```
[50]: df.shape
```

```
[50]: (20640, 10)
```

```
[51]: df.describe()
```

```
[51]: longitude latitude housing_median_age total_rooms \
count 20640.000000 20640.000000 20640.000000 20640.000000
mean -119.569704 35.631861 28.639486 2635.763081
std 2.003532 2.135952 12.585558 2181.615252
```

min	-124.350000	32.540000	1.000000	2.000000
25%	-121.800000	33.930000	18.000000	1447.750000
50%	-118.490000	34.260000	29.000000	2127.000000
75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
[52]: df.isnull().sum()
```

```
[52]: longitude      0
latitude           0
housing_median_age  0
total_rooms        0
total_bedrooms     207
population         0
households         0
median_income      0
median_house_value  0
ocean_proximity    0
dtype: int64
```

```
[53]: df.dropna(inplace=True)
```

```
[54]: df.isnull().sum()
```

```
[54]: longitude      0
latitude           0
housing_median_age  0
```

```
total_rooms      0
total_bedrooms   0
population       0
households       0
median_income    0
median_house_value 0
ocean_proximity  0
dtype: int64
```

```
[55]: numerical_cols = df.select_dtypes(include=["float64", "int64", "number"])
      categorical_cols = df.select_dtypes(include=["object"])
```

```
[56]: numerical_cols_names = numerical_cols.columns
      categorical_cols_names = categorical_cols.columns
```

```
[57]: mean_values = np.mean(numerical_cols, axis=0)
      median_values = np.median(numerical_cols, axis=0)
      std_dev = np.std(numerical_cols, axis=0)
      quartiles = np.percentile(numerical_cols, [25, 50, 75], axis=0)
```

```
[58]: print(f"{mean_values=}\n\n{median_values=}\n\n{std_dev=}\n\n{quartiles=}")
```

```
mean_values=longitude      -119.570689
latitude                   35.633221
housing_median_age         28.633094
total_rooms                2636.504233
total_bedrooms             537.870553
population                 1424.946949
households                 499.433465
median_income              3.871162
median_house_value         206864.413155
dtype: float64
```

```
median_values=array([-1.1849e+02,  3.4260e+01,  2.9000e+01,  2.1270e+03,
 4.3500e+02,
 1.1660e+03,  4.0900e+02,  3.5365e+00,  1.7970e+05])
```

```
std_dev=longitude          2.003529
latitude                   2.136295
housing_median_age         12.591497
total_rooms                2185.216092
total_bedrooms             421.374759
population                 1133.180760
households                 382.289871
median_income              1.899245
```


```
median_house_value    115432.842328
dtype: float64
```

```
quartiles=array([[ -1.2180e+02,  3.3930e+01,  1.8000e+01,  1.4500e+03,
 2.9600e+02,
 7.8700e+02,  2.8000e+02,  2.5637e+00,  1.1950e+05],
 [-1.1849e+02,  3.4260e+01,  2.9000e+01,  2.1270e+03,  4.3500e+02,
 1.1660e+03,  4.0900e+02,  3.5365e+00,  1.7970e+05],
 [-1.1801e+02,  3.7720e+01,  3.7000e+01,  3.1430e+03,  6.4700e+02,
 1.7220e+03,  6.0400e+02,  4.7440e+00,  2.6470e+05]])
```

```
[59]: corr_matrix = numerical_cols.corr()

plt.figure(figsize=(15, 10))
sns.heatmap(corr_matrix, cmap="coolwarm", annot=True, linewidths=0.5)
```

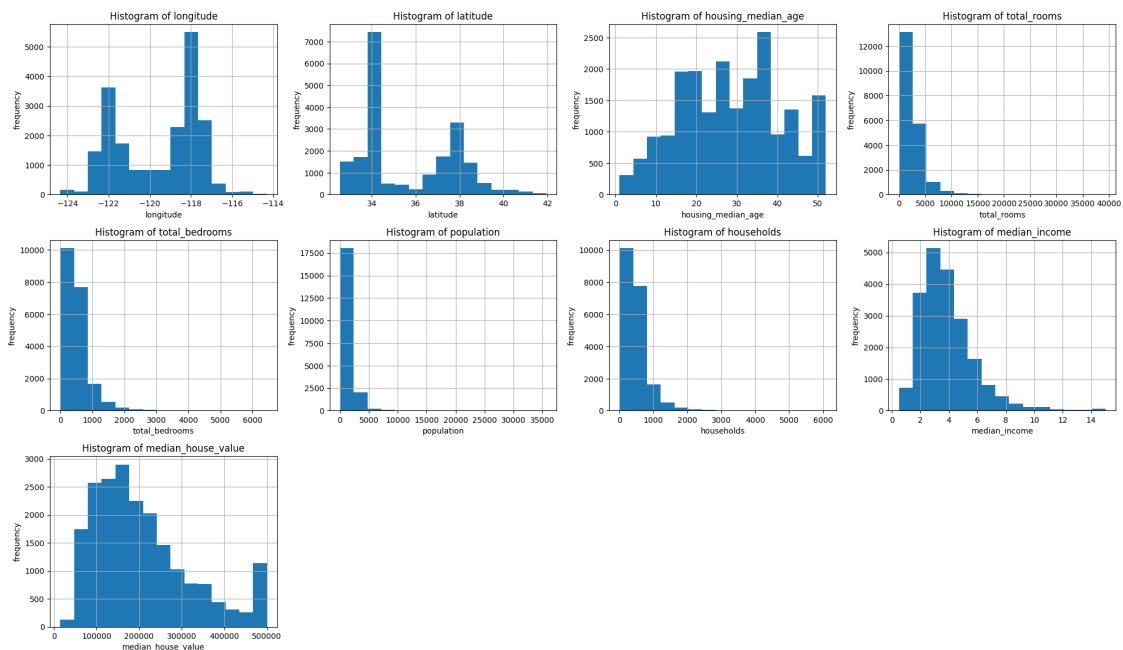
```
[59]: <Axes: >
```



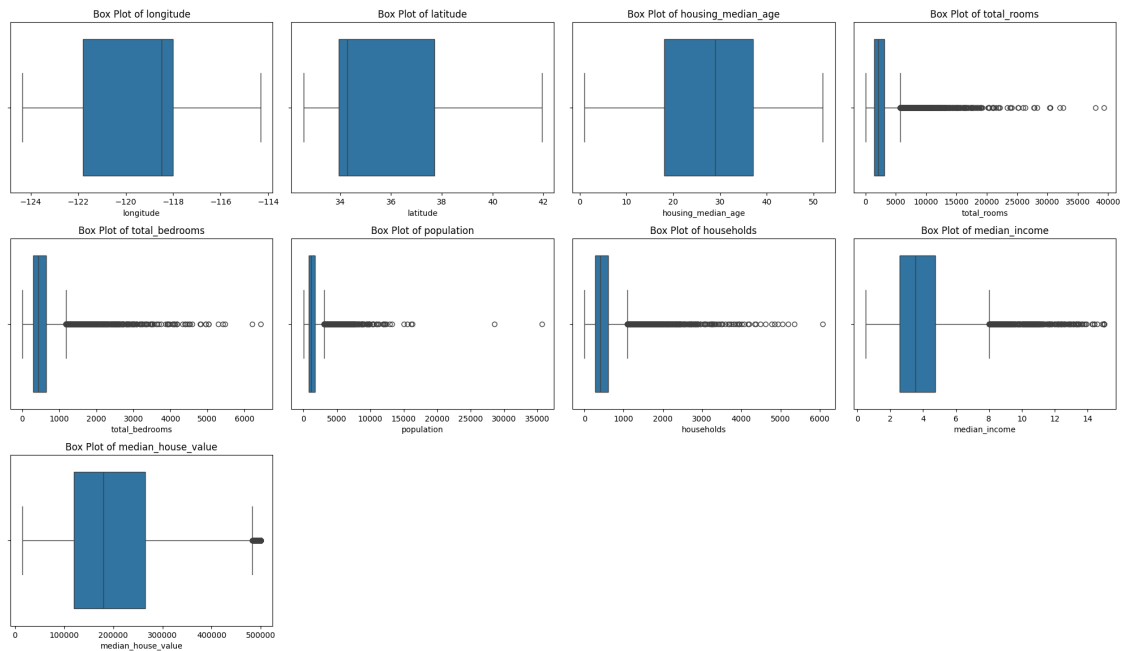
```
[60]: target_corr = corr_matrix["median_income"].sort_values(ascending=False)
target_corr[1:6]
```

```
[60]: median_house_value    0.688355
      total_rooms          0.197882
      households           0.013434
      population           0.005087
      total_bedrooms       -0.007723
      Name: median_income, dtype: float64
```

```
[61]: plt.figure(figsize=(20, 15))
      for i, col in enumerate(numerical_cols_names):
          plt.subplot(4, 4, i+1)
          df[col].hist(bins=15)
          plt.xlabel(col)
          plt.ylabel("frequency")
          plt.title(f"Histogram of {col}")
      plt.tight_layout()
      plt.show()
```

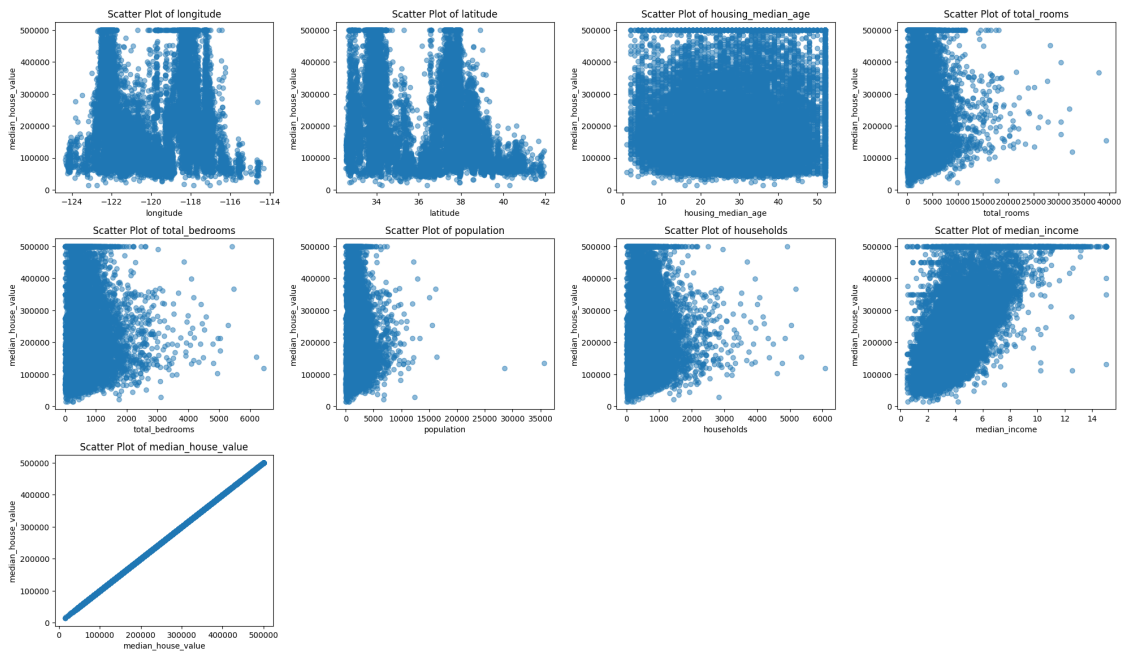


```
[62]: plt.figure(figsize=(20, 15))
      for i, col in enumerate(numerical_cols_names):
          plt.subplot(4, 4, i+1)
          sns.boxplot(data=df, x=col)
          plt.xlabel(col)
          plt.title(f"Box Plot of {col}")
      plt.tight_layout()
      plt.show()
```



```
[63]: plt.figure(figsize=(20, 15))
target = "median_house_value"

for i, col in enumerate(numerical_cols_names):
    plt.subplot(4, 4, i+1)
    plt.scatter(df[col], df[target], alpha=0.5)
    plt.xlabel(col)
    plt.ylabel(target)
    plt.title(f"Scatter Plot of {col}")
plt.tight_layout()
plt.show()
```



```
[64]: df['income_per_room'] = df['median_income'] / df['total_rooms']
df['average_household_per_population'] = df['households'] / df['population']
print(df['income_per_room'])
print(df['average_household_per_population'])
```

```
0      0.009460
1      0.001169
2      0.004947
3      0.004429
4      0.002364
```

...

```
20635    0.000937
20636    0.003668
20637    0.000754
20638    0.001004
20639    0.000858
```

Name: income\_per\_room, Length: 20433, dtype: float64

```
0      0.391304
1      0.473969
2      0.356855
3      0.392473
4      0.458407
```

...

```
20635    0.390533
20636    0.320225
20637    0.429990
```



```
20638    0.470985
20639    0.382120
Name: average_household_per_population, Length: 20433, dtype: float64
```

```
[65]: def find_outliers(df, column):
        Q1 = np.percentile(df[column], 25)
        Q3 = np.percentile(df[column], 75)
        IQR = Q3 - Q1
        lower_bound = Q1 - (1.5 * IQR)
        upper_bound = Q3 + (1.5 * IQR)

        return (df[column] < lower_bound) | (df[column] > upper_bound)
```

```
[66]: df_outliers_removed = df.copy()
        for col in numerical_cols_names:
            outliers = find_outliers(df, col)
            print(col, outliers.sum())
            df_outliers_removed = df_outliers_removed[~outliers]
            # df_outliers_removed.loc[outliers, col] = df_outliers_removed[col].median()
```

```
longitude 0
latitude 0
housing_median_age 0
total_rooms 1290
total_bedrooms 1271
population 1190
households 1210
median_income 670
median_house_value 1064
```

```
C:\Users\Pranesh\AppData\Local\Temp\ipykernel_1776\76152672.py:5: UserWarning:
Boolean Series key will be reindexed to match DataFrame index.
```

```
df_outliers_removed = df_outliers_removed[~outliers]
```

```
C:\Users\Pranesh\AppData\Local\Temp\ipykernel_1776\76152672.py:5: UserWarning:
Boolean Series key will be reindexed to match DataFrame index.
```

```
df_outliers_removed = df_outliers_removed[~outliers]
```

```
C:\Users\Pranesh\AppData\Local\Temp\ipykernel_1776\76152672.py:5: UserWarning:
Boolean Series key will be reindexed to match DataFrame index.
```

```
df_outliers_removed = df_outliers_removed[~outliers]
```

```
C:\Users\Pranesh\AppData\Local\Temp\ipykernel_1776\76152672.py:5: UserWarning:
Boolean Series key will be reindexed to match DataFrame index.
```

```
df_outliers_removed = df_outliers_removed[~outliers]
```

```
C:\Users\Pranesh\AppData\Local\Temp\ipykernel_1776\76152672.py:5: UserWarning:
Boolean Series key will be reindexed to match DataFrame index.
```

```
df_outliers_removed = df_outliers_removed[~outliers]
```

```
[67]: for col in numerical_cols_names:
        outliers = find_outliers(df_outliers_removed, col)
```

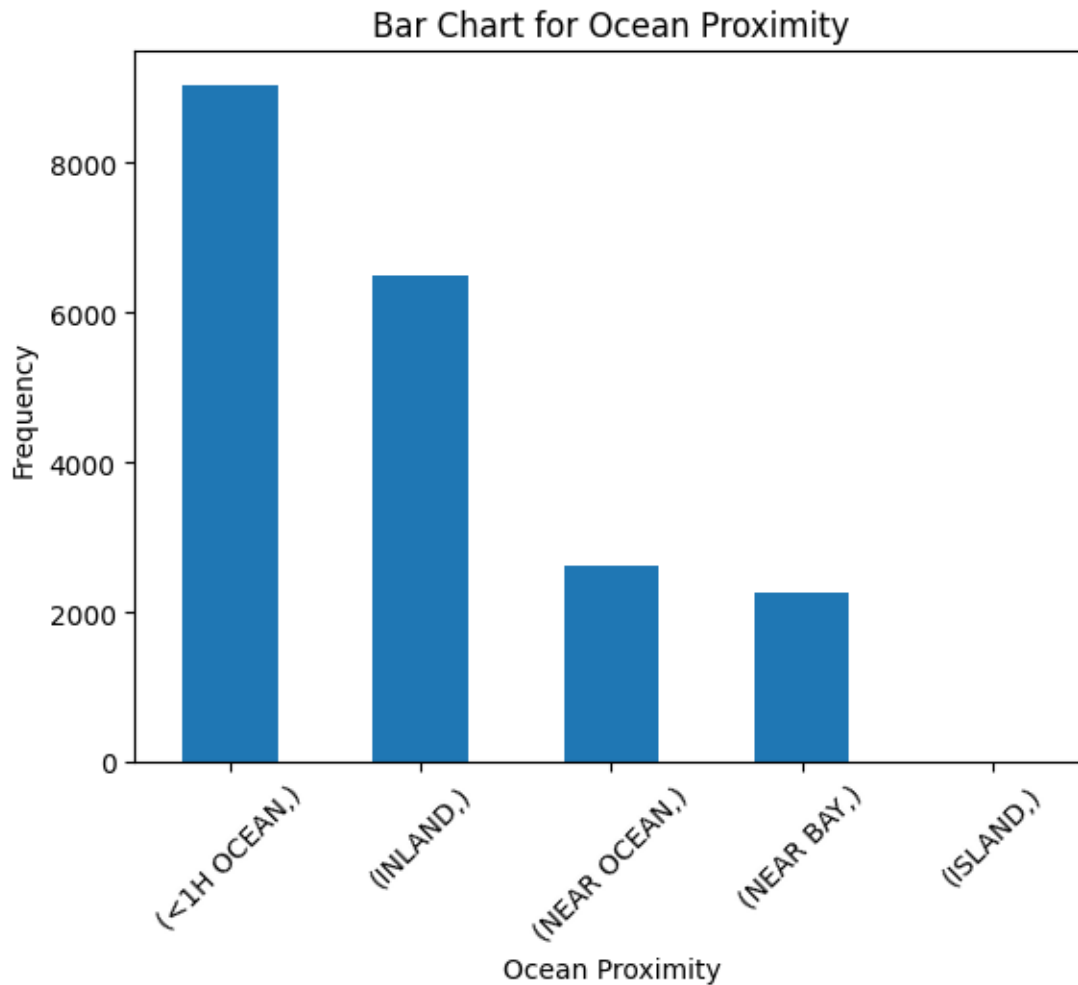
```
print(col, outliers.sum())
```

```
longitude 0  
latitude 0  
housing_median_age 0  
total_rooms 322  
total_bedrooms 273  
population 288  
households 272  
median_income 175  
median_house_value 261
```

```
[68]: df = df_outliers_removed
```

```
[69]: categorical_cols.value_counts().plot(kind='bar')  
plt.xlabel("Ocean Proximity")  
plt.ylabel("Frequency")  
plt.title("Bar Chart for Ocean Proximity")  
plt.xticks(rotation=45)
```

```
[69]: (array([0, 1, 2, 3, 4]),  
      [Text(0, 0, '(<1H OCEAN,)',  
          Text(1, 0, '(INLAND,)',  
          Text(2, 0, '(NEAR OCEAN,)',  
          Text(3, 0, '(NEAR BAY,)',  
          Text(4, 0, '(ISLAND,))']])
```



```
[70]: for col in numerical_cols_names:
      print(f"Mean of {col}: ", np.mean(df[col], axis=0))
      print(f"Median of {col}: ", np.median(df[col], axis=0))
      print(f"Mode of {col}: ", df[col].mode()[0])
      print("\n")
```

```
Mean of longitude: -119.60577262819778
Median of longitude: -118.61
Mode of longitude: -118.31
```

```
Mean of latitude: 35.69786164965011
Median of latitude: 34.31
Mode of latitude: 34.08
```

Mean of housing\_median\_age: 29.489216473557416  
Median of housing\_median\_age: 30.0  
Mode of housing\_median\_age: 52.0

Mean of total\_rooms: 2144.865205919468  
Median of total\_rooms: 1979.0  
Mode of total\_rooms: 1527.0

Mean of total\_bedrooms: 445.4726970287943  
Median of total\_bedrooms: 411.0  
Mode of total\_bedrooms: 280.0

Mean of population: 1197.1082941378916  
Median of population: 1111.0  
Mode of population: 1227.0

Mean of households: 416.80612596076634  
Median of households: 387.0  
Mode of households: 306.0

Mean of median\_income: 3.575592640816795  
Median of median\_income: 3.3906  
Mode of median\_income: 3.125

Mean of median\_house\_value: 187056.1888264311  
Median of median\_house\_value: 170100.0  
Mode of median\_house\_value: 137500.0

```
[71]: for col in numerical_cols_names:
        df[col] = (df[col] - df[col].min()) / (df[col].max() - df[col].min())

        # or
        # scaler = StandardScaler()

        # df[numerical_cols] = scaler.fit_transform(df[numerical_cols])

        # df.head()
```

```
[72]: df
```

```
[72]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
2	0.213996	0.564293	1.000000	0.258241	0.160547	
3	0.212982	0.564293	1.000000	0.224220	0.198975	
4	0.212982	0.564293	1.000000	0.286445	0.237404	
5	0.212982	0.564293	1.000000	0.161643	0.180188	
6	0.212982	0.563231	1.000000	0.446501	0.415884	
...	...	...	...	...	...	
20635	0.330629	0.737513	0.470588	0.293143	0.317677	
20636	0.318458	0.738576	0.333333	0.122510	0.126388	
20637	0.317444	0.732200	0.313725	0.396968	0.412468	
20638	0.307302	0.732200	0.333333	0.327516	0.347566	
20639	0.315416	0.725824	0.294118	0.490569	0.524338	

	population	households	median_income	median_house_value	\
2	0.157962	0.160846	0.899633	0.721533	
3	0.177828	0.199449	0.684719	0.698417	
4	0.180070	0.236213	0.445496	0.700343	
5	0.131368	0.175551	0.470871	0.545164	
6	0.349567	0.470588	0.420587	0.608306	
...	...	...	...	...	
20635	0.269785	0.301471	0.141172	0.135062	
20636	0.113105	0.102941	0.273837	0.132921	
20637	0.321692	0.396140	0.159770	0.165456	
20638	0.236463	0.318934	0.182030	0.149188	
20639	0.443448	0.485294	0.251444	0.159248	

	ocean_proximity	income_per_room	average_household_per_population
2	NEAR BAY	0.004947	0.356855
3	NEAR BAY	0.004429	0.392473
4	NEAR BAY	0.002364	0.458407
5	NEAR BAY	0.004393	0.467312
6	NEAR BAY	0.001443	0.469835
...	...	...	...
20635	INLAND	0.000937	0.390533
20636	INLAND	0.003668	0.320225
20637	INLAND	0.000754	0.429990
20638	INLAND	0.001004	0.470985
20639	INLAND	0.000858	0.382120

[17434 rows x 12 columns]

```
[73]: target_col = 'median_income'
```


```
[74]: plt.figure(figsize=(15, 10))

plt.subplot(2, 2, 1)
df[target_col].hist(bins=15)
```

```
plt.xlabel(target_col)
plt.ylabel("Frequency")
plt.title(f'Histogram of {target_col}')
```


```
plt.subplot(2, 2, 2)
sns.histplot(df[target_col], bins=15, kde=True, color='skyblue')
plt.xlabel(target_col)
plt.ylabel("Frequency")
plt.title(f'Histogram of {target_col}')
```

[74]: Text(0.5, 1.0, 'Histogram of median\_income')



[75]: plt.subplot(2, 1, 1)

```
stats.probplot(df[target_col], dist="norm", plot=plt)
plt.title(f'QQ Plot of {target_col}')
plt.tight_layout()
plt.show()
```



```
[76]: shapiro_stat, shapiro_p = stats.shapiro(df[target_col])
```

```
C:\Users\Pranesh\AppData\Local\Programs\Python\Python39\lib\site-  
packages\scipy\stats\_morestats.py:1882: UserWarning: p-value may not be  
accurate for N > 5000.
```

```
warnings.warn("p-value may not be accurate for N > 5000.")
```

```
[77]: kstest_stat, kstest_p = stats.kstest(df[target_col], "norm",  
    ↪ args=(df[target_col].mean(), df[target_col].std()))
```

```
[78]: print(shapiro_p, kstest_p)
```

```
0.0 5.49876265719009e-48
```

```
[79]: if shapiro_p > 0.05 and kstest_p > 0.05:  
    print("Data appears to be normally distributed")  
else:  
    print("Data does not appear to be normally distributed")
```

```
Data does not appear to be normally distributed
```


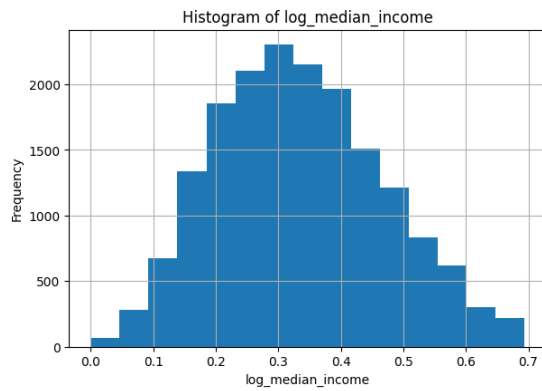
```
[80]: df[f'log_{target_col}'] = np.log1p(df[target_col])
```

```
[81]: plt.figure(figsize=(15, 10))  
  
plt.subplot(2, 2, 1)  
df[f'log_{target_col}'].hist(bins=15)  
plt.xlabel(f'log_{target_col}')  
plt.ylabel("Frequency")  
plt.title(f'Histogram of log_{target_col}')
```

```
plt.subplot(2, 2, 2)  
sns.histplot(df[f'log_{target_col}'], bins=15, kde=True, color='skyblue')  
plt.xlabel(f'log_{target_col}')
```


```
plt.ylabel("Frequency")  
plt.title(f'Histogram of log_{target_col}')
```

```
[81]: Text(0.5, 1.0, 'Histogram of log_median_income')
```



```
[82]: plt.subplot(2, 1, 1)

stats.probplot(df[f'log_{target_col}'], dist="norm", plot=plt)
plt.title(f'QQ Plot of log_{target_col}')
plt.tight_layout()
plt.show()
```



```
[83]: help(stats.boxcox)
```

Help on function boxcox in module scipy.stats.\_morestats:

```
boxcox(x, lmbda=None, alpha=None, optimizer=None)
```

Return a dataset transformed by a Box-Cox power transformation.

Parameters

-----

x : ndarray



Input array to be transformed.

If ``lambda`` is not None, this is an alias of

``scipy.special.boxcox``.

Returns nan if ``x` < 0``; returns -inf if ``x` == 0` and `lambda < 0``.

If ``lambda`` is None, array must be positive, 1-dimensional, and non-constant.

`lambda` : scalar, optional

If ``lambda`` is None (default), find the value of ``lambda`` that maximizes the log-likelihood function and return it as the second output argument.

If ``lambda`` is not None, do the transformation for that value.

`alpha` : float, optional

If ``lambda`` is None and ``alpha`` is not None (default), return the ``100 * (1-alpha)%`` confidence interval for ``lambda`` as the third output argument. Must be between 0.0 and 1.0.

If ``lambda`` is not None, ``alpha`` is ignored.

`optimizer` : callable, optional

If ``lambda`` is None, ``optimizer`` is the scalar optimizer used to find the value of ``lambda`` that minimizes the negative log-likelihood function. ``optimizer`` is a callable that accepts one argument:

`fun` : callable

The objective function, which evaluates the negative log-likelihood function at a provided value of ``lambda``

and returns an object, such as an instance of

``scipy.optimize.OptimizeResult``, which holds the optimal value of ``lambda`` in an attribute ``x``.

See the example in ``boxcox_normmax`` or the documentation of ``scipy.optimize.minimize_scalar`` for more information.

If ``lambda`` is not None, ``optimizer`` is ignored.

Returns

-----

`boxcox` : ndarray

Box-Cox power transformed array.

`maxlog` : float, optional

If the ``lambda`` parameter is None, the second returned argument is the ``lambda`` that maximizes the log-likelihood function.

`(min_ci, max_ci)` : tuple of float, optional

If ``lambda`` parameter is None and ``alpha`` is not None, this returned tuple of floats represents the minimum and maximum confidence limits given ``alpha``.

See Also

-----

`probplot`, `boxcox_normplot`, `boxcox_normmax`, `boxcox_llf`

Notes

-----

The Box-Cox transform is given by::

$$y = \begin{cases} (x^{\lambda} - 1) / \lambda, & \text{for } \lambda \neq 0 \\ \log(x), & \text{for } \lambda = 0 \end{cases}$$

``boxcox`` requires the input data to be positive. Sometimes a Box-Cox transformation provides a shift parameter to achieve this; ``boxcox`` does not. Such a shift parameter is equivalent to adding a positive constant to ``x`` before calling ``boxcox``.

The confidence limits returned when ``alpha`` is provided give the interval where:

.. math::

$$llf(\hat{\lambda}) - llf(\lambda) < \frac{1}{2} \chi^2(1 - \alpha, 1),$$

with ``llf`` the log-likelihood function and `:math:`\chi^2`` the chi-squared function.

References

-----

G.E.P. Box and D.R. Cox, "An Analysis of Transformations", Journal of the Royal Statistical Society B, 26, 211-252 (1964).

Examples

-----

```
>>> from scipy import stats
>>> import matplotlib.pyplot as plt
```

We generate some random variates from a non-normal distribution and make a probability plot for it, to show it is non-normal in the tails:

```
>>> fig = plt.figure()
>>> ax1 = fig.add_subplot(211)
>>> x = stats.loggamma.rvs(5, size=500) + 5
>>> prob = stats.probplot(x, dist=stats.norm, plot=ax1)
>>> ax1.set_xlabel('')
```

```
>>> ax1.set_title('Probplot against normal distribution')
```

We now use `boxcox` to transform the data so it's closest to normal:

```
>>> ax2 = fig.add_subplot(212)
>>> xt, _ = stats.boxcox(x)
>>> prob = stats.probplot(xt, dist=stats.norm, plot=ax2)
>>> ax2.set_title('Probplot after Box-Cox transformation')

>>> plt.show()
```

```
[84]: df.columns
```

```
[84]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
        'total_bedrooms', 'population', 'households', 'median_income',
        'median_house_value', 'ocean_proximity', 'income_per_room',
        'average_household_per_population', 'log_median_income'],
        dtype='object')
```

```
[85]: X = df[['median_income']] # independent (input); double brackets is to convert
      ↪ it into a data frame (needs to be 2D in fitting)
      Y = df['median_house_value'] # dependent (output)
```

```
[86]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
      ↪ random_state = 42)
```

```
[ ]:
```

```
[87]: model = LinearRegression()
      model.fit(X_train, Y_train)
```

```
[87]: LinearRegression()
```

```
[88]: intercept = model.intercept_
      slope = model.coef_[0]
      print(f"Regression Eqn: median_house_value = {slope:4f} * median_income +
      ↪ {intercept:.4f}")
```


Regression Eqn: median\_house\_value = 0.660493 \* median\_income + 0.0983

```
[89]: y_pred = model.predict(X_test)
```

```
[90]: plt.figure(figsize=(15, 10))
      plt.scatter(X_test, Y_test, color='blue', label='Actual Values', alpha=0.5) #
      ↪ for give x, what is the actual y
```

```
plt.plot(X_test, y_pred, color="red", label="Regression Line", linewidth=2) #
    ↪ for given x, what is the predicted y
plt.title('Regression Line')
plt.xlabel('Median Income')
plt.ylabel('Median House Value')
plt.legend()
```

[90]: <matplotlib.legend.Legend at 0x246faca68e0>



```
[91]: mse = mean_squared_error(Y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(Y_test, y_pred)
r_squared = r2_score(Y_test, y_pred)

# Print the metrics
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"R-squared: {r_squared:.2f}")
```

Mean Squared Error (MSE): 0.02  
 Root Mean Squared Error (RMSE): 0.16  
 Mean Absolute Error (MAE): 0.12

R-squared: 0.38

```
[92]: plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
residuals = Y_test - y_pred
plt.scatter(y_pred, residuals, color="orange", alpha=0.5)
plt.axhline(0, color="red", linestyle="--")
plt.title('Predicted Values vs Residuals')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')

plt.tight_layout()
plt.show()
```

